Empirical Economics Cheat-Sheet

Asymptotics

Modes of Convergence

Let $\{X_n\}_{n\geq 1}$ and X be random variables on $(\Omega, \mathcal{F}, P)$.

Almost Sure Convergence: $X_n \xrightarrow{a.s.} X$ if $P(\{\omega : X_n(\omega) \to X(\omega)\}) = 1$.

Convergence in Probability: $X_n \xrightarrow{p} X$ if $\forall \epsilon > 0$: $P(|X_n - X| > \epsilon) \to 0$.

Convergence in r-th Mean: $X_n \xrightarrow{r} X$ if $\mathbb{E}(|X_n - X|^r) \to 0$.

Convergence in Distribution: $X_n \xrightarrow{d} X$ if $F_{X_n}(x) \to F_X(x)$ for all x where F_X is continuous.

Implications between modes

$$X_n \xrightarrow{a.s.} X \implies X_n \xrightarrow{p} X \implies X_n \xrightarrow{d} X$$
$$X_n \xrightarrow{r} X \implies X_n \xrightarrow{p} X \implies X_n \xrightarrow{d} X$$
$$X_n \xrightarrow{r} X \implies X_n \xrightarrow{s} X \text{ for } s \leq r$$

None of the reverse implications hold in general. Exception: $X_n \xrightarrow{d} c$ (constant) $\implies X_n \xrightarrow{p} c$.

Continuous Mapping Theorem

Let $g : \mathbb{R}^k \to \mathbb{R}^n$ be continuous on $S \subset \mathbb{R}^k$ with $P(X \in S) = 1$. Then:

$$X_n \xrightarrow{a.s.} X \implies g(X_n) \xrightarrow{a.s.} g(X)$$
$$X_n \xrightarrow{p} X \implies g(X_n) \xrightarrow{p} g(X)$$
$$X_n \xrightarrow{d} X \implies g(X_n) \xrightarrow{d} g(X)$$

Does NOT hold for $\xrightarrow{r}$. Counterexample: $X_n = n$ w.p. $1/n^2$, 0 otherwise. $g(x) = x^2$: $\mathbb{E}|X_n| \to 0$ but $\mathbb{E}|X_n^2| = 1$.
Important: Need $P(X \in S) = 1$. E.g. $g(x, y) = x/y$ continuous on $S = \mathbb{R}^2 \setminus \{(x, 0)\}$; need $c \neq 0$ for $X_n/Y_n \xrightarrow{d} X/c$.

Slutsky's Theorem

If $X_n \xrightarrow{d} X$ and $Y_n \xrightarrow{p} c$ (constant), then:

$$X_n + Y_n \xrightarrow{d} X + c$$
$$X_n Y_n \xrightarrow{d} cX$$
$$X_n/Y_n \xrightarrow{d} X/c \quad (c \neq 0)$$

Critical: Y_n must converge to a *constant*. If $Y_n \xrightarrow{d} Y$ (non-degenerate), Slutsky does not apply.

Weak Law of Large Numbers

If $\{X_i\}_{i\geq 1}$ iid with $\mathbb{E}(X_i) = \mu$, $\text{Var}(X_i) = \sigma^2 < \infty$, then $\bar{X}_n \xrightarrow{p} \mu$.
Proof (Chebyshev): $P(|\bar{X}_n - \mu| > \epsilon) \leq \frac{\text{Var}(\bar{X}_n)}{\epsilon^2} = \frac{\sigma^2}{n\epsilon^2} \to 0$.

Central Limit Theorem

If $\{X_i\}_{i\geq 1}$ iid with $\mathbb{E}(X_i) = \mu$, $\text{Var}(X_i) = \Sigma$ (finite), then:

$$\sqrt{n}(\bar{X}_n - \mu) \xrightarrow{d} N(0, \Sigma).$$

Stochastic Order Notation

$X_n = O_p(1)$: $\{X_n\}$ is bounded in probability, i.e. $\forall \epsilon > 0$, $\exists M$ s.t. $\sup_n P(|X_n| > M) < \epsilon$.
$X_n = o_p(1)$: $X_n \xrightarrow{p} 0$.
Key rules: $O_p(1) \cdot o_p(1) = o_p(1)$; $\quad O_p(1) + O_p(1) = O_p(1)$.
$X_n \xrightarrow{d} X \implies X_n = O_p(1)$.
If $\sqrt{n}(X_n - c) \xrightarrow{d} X$, then $X_n \xrightarrow{p} c$ and $\sqrt{n}(X_n - c) = O_p(1)$.

Delta Method

Delta Method (General)

Let $\{X_n\}_{n\geq 1}$ be $(K \times 1)$ random vectors with $n^r(X_n - c) \xrightarrow{d} X$ for some $r > 0$ and constant c. Let $g : \mathbb{R}^k \to \mathbb{R}^d$ be differentiable at c with Jacobian $Dg(c)$. Then:

$$n^r(g(X_n) - g(c)) \xrightarrow{d} Dg(c)\, X.$$

If $X \sim N(0, \Sigma)$:

$$n^r(g(X_n) - g(c)) \xrightarrow{d} N\big(0,\, Dg(c)\, \Sigma\, Dg(c)'\big)\, .$$

Proof Sketch

By Taylor: $g(x) = g(c) + Dg(c)(x - c) + h_1(x)(x - c)$ with $h_1(c) = 0$. Then:

$$n^r(g(X_n) - g(c)) = Dg(c)\, n^r(X_n - c) + h_1(X_n)\, n^r(X_n - c).$$

Since $X_n \xrightarrow{p} c$, CMT gives $h_1(X_n) \xrightarrow{p} 0$, and $n^r(X_n - c) = O_p(1)$, so the remainder is $o_p(1) \cdot O_p(1) = o_p(1)$.

Second-Order Delta Method

If $g'(c) = 0$ and $g''(c)$ exists (scalar case), then:

$$n^{2r}(g(X_n) - g(c)) \xrightarrow{d} \frac{g''(c)}{2} X^2.$$

Use when first-order term vanishes (degenerate limit).

Application: Sample Variance

Let X_i iid with $\mathbb{E}(X_i) = \mu$, $\text{Var}(X_i) = \sigma^2$, $\mathbb{E}(X_i - \mu)^4 = \kappa$. Let $S_n^2 = \frac{1}{n}\sum_{i=1}^n (X_i - \bar{X})^2$. Using the Delta Method on $g(\mu, m_2) = m_2 - \mu^2$ applied to the sample moments $(\bar{X}_n, \frac{1}{n}\sum X_i^2)$:

$$\sqrt{n}(S_n^2 - \sigma^2) \xrightarrow{d} N(0, \kappa - \sigma^4).$$

Estimation

Definitions

Given a sample $\{X_i\}_{i=1}^n$ from distribution F, a **statistic** is a function $T_n : (X_1, \ldots, X_n) \to V$. An **estimator** is a statistic used to learn about some feature $\theta(F)$.

Finite Sample Properties

Bias: $\text{Bias}(\hat{\theta}_n) = \mathbb{E}(\hat{\theta}_n) - \theta$. Unbiased if $\mathbb{E}(\hat{\theta}_n) = \theta$.
Mean Squared Error:

$$\text{MSE}(\hat{\theta}_n) = \mathbb{E}\Big[(\hat{\theta}_n - \theta)^2\Big] = \text{Var}(\hat{\theta}_n) + \text{Bias}(\hat{\theta}_n)^2.$$

Large Sample Properties

Consistency: $\hat{\theta}_n \xrightarrow{p} \theta$ (or $\hat{\theta}_n \xrightarrow{a.s.} \theta$).
Asymptotic Normality: $\sqrt{n}(\hat{\theta}_n - \theta) \xrightarrow{d} N(0, V)$ for some V.
Asymptotic Efficiency: An estimator achieving the smallest possible asymptotic variance among regular estimators (e.g. MLE under regularity conditions).

Method of Moments

Sample Analogue Principle: Replace population moments with sample moments.
If θ satisfies $\mathbb{E}(m(X, \theta)) = 0$ for moment function m, the MoM estimator solves:

$$\frac{1}{n}\sum_{i=1}^n m(X_i, \hat{\theta}_n) = 0.$$

Consistency follows from SLLN + CMT if identification holds.

Maximum Likelihood Estimation

Setup

Let $\{X_i\}_{i=1}^n$ iid with density f_{θ_0} for some $\theta_0 \in \Theta \subset \mathbb{R}^d$.
Likelihood: $\ell_n(\theta) = \prod_{i=1}^n f_\theta(X_i)$.
Log-likelihood: $L_n(\theta) = \frac{1}{n}\sum_{i=1}^n \ln f_\theta(X_i)$.
MLE: $\hat{\theta}_n \in \arg\max_{\theta\in\Theta} L_n(\theta)$.

Why MLE is Consistent

Under regularity conditions, $L_n(\theta) \xrightarrow{p} L(\theta) := \mathbb{E}(\ln f_\theta(X))$.
Key fact: θ_0 uniquely maximizes $L(\theta)$.
Proof: Let $M(\theta) = L(\theta) - L(\theta_0) = \mathbb{E}\Big[\ln \frac{f_\theta(X)}{f_{\theta_0}(X)}\Big]$. By Jensen's inequality:

$$M(\theta) \leq \ln \mathbb{E}\Big[\frac{f_\theta(X)}{f_{\theta_0}(X)}\Big] = \ln \int \frac{f_\theta(x)}{f_{\theta_0}(x)} f_{\theta_0}(x)\, dx = \ln 1 = 0,$$

with equality iff $f_\theta(X)/f_{\theta_0}(X) = c$ a.s. Ruled out for $\theta \neq \theta_0$ by assumption.

Asymptotic Distribution

Under regularity conditions:

$$\sqrt{n}(\hat{\theta}_{\text{MLE}} - \theta_0) \xrightarrow{d} N\big(0,\ I(\theta_0)^{-1}\big)\, ,$$

where $I(\theta_0) = -\mathbb{E}\Big[\frac{\partial^2}{\partial\theta^2} \ln f_{\theta_0}(Y|X)\Big]$ is the **Fisher Information**. This variance is optimal: no regular estimator can achieve a smaller asymptotic variance.

Example: Normal

$X_i \sim N(\mu, \sigma^2)$, $\theta = (\mu, \sigma^2)$ unknown.

$$L_n(\theta) = -\frac{1}{2}\ln(2\pi) - \frac{1}{2}\ln\sigma^2 - \frac{1}{2n\sigma^2}\sum_{i=1}^{n}(X_i - \mu)^2.$$

FOCs yield $\hat{\mu} = \bar{X}_n$, $\hat{\sigma}^2 = \frac{1}{n}\sum(X_i - \bar{X}_n)^2$.

Example: Bernoulli

$X_i \sim$ Bernoulli(θ). Log-likelihood:
$L_n = \bar{X}_n \ln\theta + (1-\bar{X}_n)\ln(1-\theta)$. FOC: $\hat{\theta} = \bar{X}_n$. Log-likelihood is concave, so FOC suffices.

Conditional MLE

If $Y|X$ has conditional density $f_\theta(y|x)$: $\ell_n(\theta) = \prod_i f_\theta(Y_i|X_i)$. Maximising this is equivalent to OLS when $Y|X \sim N(X'\beta, \sigma^2)$.

Normal Regression as CMLE

If $Y|X \sim N(\beta_0 + \beta_1 X, \sigma^2)$, CMLE of (β_0, β_1) minimizes $\sum(Y_i - \beta_0 - \beta_1 X_i)^2$. These are the OLS estimators.

OLS: Setup & Projections

Linear Model

$$y_i = x_i'\beta + u_i, \quad \mathbb{E}(x_i u_i) = 0.$$

$x_i \in \mathbb{R}^{k+1}$ with $x_{i0} = 1$. "Linear" means linear in parameters β_j. The error u_i contains unobserved determinants of y_i.

Identification

Assume $\mathbb{E}(xu) = 0$ and no perfect collinearity (no $a \neq 0$ with $P(a'x = 0) = 1$).
$\mathbb{E}(xx')$ invertible $\iff$ no perfect collinearity. Then:

$$\beta = \mathbb{E}(xx')^{-1}\mathbb{E}(xy).$$

Proof ($\mathbb{E}(xx')$ invertible $\iff$ no collinearity):
($\Rightarrow$) If $P(x'a = 0) = 1$ for $a \neq 0$, then $\mathbb{E}(xx')a = \mathbb{E}(x \cdot x'a) = 0$, not invertible.
($\Leftarrow$) No collinearity $\implies$ $c'\mathbb{E}(xx')c = \mathbb{E}[(x'c)^2] > 0\ \forall c \neq 0$, so $\mathbb{E}(xx')$ is positive definite.

OLS Estimator

Given iid sample $\{y_i, x_i\}_{i=1}^n$. Unique OLS estimator (when $X'X$ invertible):

$$\hat{\beta}_n = \left(\frac{1}{n}\sum x_i x_i'\right)^{-1}\frac{1}{n}\sum x_i y_i = (X'X)^{-1}X'Y.$$

Equivalently solves $\min_b \|Y - Xb\|^2$. FOC: $X'\hat{U} = 0$.

Projection Matrix

$P_X = X(X'X)^{-1}X'$: projects onto column space $\mathcal{S}(X)$.
$M_X = I_n - P_X$: residual maker.
Properties:

- $P_X = P_X'$, $M_X = M_X'$ (symmetric)
- $P_X^2 = P_X$, $M_X^2 = M_X$ (idempotent)
- $P_X M_X = M_X P_X = 0$
- $P_X X = X$, $M_X X = 0$
- For any Y: $Y = P_X Y + M_X Y = \hat{Y} + \hat{U}$

Projection Theorem

Let $\mathcal{S}$ be a nonempty subspace of $\mathbb{R}^n$. There exists a unique $\hat{y} \in \mathcal{S}$ minimizing $\|y - \hat{y}\|$. Necessary and sufficient: $y - \hat{y}$ is orthogonal to every vector in $\mathcal{S}$.
Applying to $\mathcal{S} = \mathcal{S}(X)$: the condition $X'(Y - \hat{Y}) = 0$ yields $\hat{Y} = P_X Y$.

Frisch-Waugh-Lovell

Partition $Y = X_1\beta_1 + X_2\beta_2 + U$. Then:

$$\hat{\beta}_2 = (X_2' M_{X_1} X_2)^{-1} X_2' M_{X_1} Y.$$

i.e., $\hat{\beta}_2$ is obtained by regressing the residuals of Y on X_1 onto the residuals of X_2 on X_1.
Proof: $M_{X_1}Y = M_{X_1}X_2\hat{\beta}_2 + \hat{U}$, multiply by X_2':
$X_2' M_{X_1} Y = X_2' M_{X_1} X_2 \hat{\beta}_2$ since $X'\hat{U} = 0$.
Population version: $\beta_2 = \mathbb{E}(\tilde{x}_2\tilde{x}_2')^{-1}\mathbb{E}(\tilde{x}_2 y)$, where $\tilde{x}_2 = x_2 - \tilde{\gamma}x_1$ is the residual from projecting x_2 onto x_1. This holds because $\mathbb{E}(\tilde{x}_2 x_1') = 0$.

Omitted Variables Bias

If $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + u$ but we regress y on x_1 only:

$$b_1 = \beta_1 + \beta_2\frac{\mathrm{Cov}(x_1, x_2)}{\mathrm{Var}(x_1)}.$$

The bias term $\beta_2\frac{\mathrm{Cov}(x_1,x_2)}{\mathrm{Var}(x_1)}$ is the effect of the omitted variable $\times$ correlation with included variable.

OLS: Finite Sample Properties

Assumptions

$Y = X\beta + U$ with $\mathbb{E}(U|X) = 0$ (equivalently $\mathbb{E}(Y|X) = X\beta$). Since (u_i, x_i) independent of x_j for $j \neq i$: $\mathbb{E}(u_i|x_1, \ldots, x_n) = 0$.

Unbiasedness

$$\mathbb{E}(\hat{\beta}_n|X) = (X'X)^{-1}X'\mathbb{E}(Y|X) = (X'X)^{-1}X'X\beta = \beta.$$

By LIE: $\mathbb{E}(\hat{\beta}_n) = \mathbb{E}[\mathbb{E}(\hat{\beta}_n|X)] = \beta$.

Variance under Homoskedasticity

Assume $\mathrm{Var}(u_i|x_i) = \sigma^2$ (homoskedastic). Then $\mathrm{Var}(U|X) = \sigma^2 I_n$ and:

$$\mathrm{Var}(\hat{\beta}_n|X) = \sigma^2(X'X)^{-1}.$$

Variance under Heteroskedasticity

If $\mathbb{E}(u_i^2|x_i) = \sigma^2(x_i)$, then $\mathrm{Var}(U|X) = \Omega$ (diagonal, varying entries):

$$\mathrm{Var}(\hat{\beta}_n|X) = (X'X)^{-1}X'\Omega X(X'X)^{-1}.$$

Gauss-Markov Theorem

Under $\mathbb{E}(U|X) = 0$ and homoskedasticity, OLS is BLUE: for any linear unbiased estimator $\tilde{\beta} = AY$ with $AX = I_{k+1}$,

$$\mathrm{Var}(\tilde{\beta}|X) - \mathrm{Var}(\hat{\beta}_n|X) = \sigma^2 CC' \succeq 0,$$

where $C = A - (X'X)^{-1}X'$.
Proof: $\mathrm{Var}(AY|X) = \sigma^2 AA'$. Let $C = A - (X'X)^{-1}X'$. Then $CX = AX - I_{k+1} = 0$, so:

$$AA' - (X'X)^{-1} = CC' + (X'X)^{-1}X'C' + CX(X'X)^{-1} = CC' \succeq 0.$$

Implication: For any $r \in \mathbb{R}^{k+1}$, $r'\hat{\beta}$ is BLUE for $r'\beta$: $\mathrm{Var}(r'\tilde{\beta}|X) - \mathrm{Var}(r'\hat{\beta}|X) = r'CC'r \geq 0$.

Unbiasedness of $\hat{\sigma}^2$

Under normality, $\hat{\sigma}^2 = \frac{\mathrm{SSR}}{n-k-1}$ is unbiased. **Proof (trace trick):**

$$\begin{aligned}\mathbb{E}[\mathrm{SSR}|X] &= \mathbb{E}[U'M_X U|X] = \mathbb{E}[\mathrm{tr}(U'M_X U)|X] \\ &= \mathbb{E}[\mathrm{tr}(M_X UU')|X] = \mathrm{tr}(M_X\,\mathbb{E}[UU'|X]) \\ &= \sigma^2\,\mathrm{tr}(M_X) = \sigma^2(n-k-1),\end{aligned}$$

since $\mathrm{tr}(M_X) = \mathrm{tr}(I_n) - \mathrm{tr}(P_X) = n - (k+1)$ (idempotent: $\mathrm{tr}(P_X) = \mathrm{tr}(X(X'X)^{-1}X') = \mathrm{tr}(I_{k+1})$).

GLS (Known Heteroskedasticity)

If $\mathrm{Var}(U|X) = \Omega$ with Ω known, pre-multiply by $\Omega^{-1/2}$: $Y^* = X^*\beta + U^*$, where $\mathrm{Var}(U^*|X) = I_n$.

$$\hat{\beta}_{\mathrm{GLS}} = (X'\Omega^{-1}X)^{-1}X'\Omega^{-1}Y.$$

This is OLS applied to the transformed model, hence BLUE by Gauss-Markov in the transformed space. Equivalently, it is BLUE in the original model.

Coefficient of Determination

$R^2 = 1 - \frac{\mathrm{SSR}}{\mathrm{TSS}} = 1 - \frac{\|M_X Y\|^2}{\|M_c Y\|^2} = \frac{\|P_X M_c Y\|^2}{\|M_c Y\|^2}$.
Where $\mathrm{TSS} = \sum(y_i - \bar{y})^2$, $\mathrm{SSR} = \sum \hat{u}_i^2$, $\mathrm{ESS} = \sum(\hat{y}_i - \bar{y})^2$.
$0 \leq R^2 \leq 1$. Always weakly increases with additional regressors.
Adjusted: $\bar{R}^2 = 1 - \frac{n-1}{n-k-1}\cdot\frac{\mathrm{SSR}}{\mathrm{TSS}} \leq R^2$.
Population: $R^2_{\mathrm{pop}} = 1 - \frac{\mathrm{Var}(u)}{\mathrm{Var}(y)}$.
High R^2 does not imply causality; low R^2 does not preclude it.

OLS: Large Sample Properties

Consistency

Under $y = x'\beta + u$, $\mathbb{E}(xu) = 0$, $\mathbb{E}(xx')$ invertible:

$$\hat{\beta}_n = \left(\frac{1}{n}\sum x_i x_i'\right)^{-1}\frac{1}{n}\sum x_i y_i \xrightarrow{a.s.} \mathbb{E}(xx')^{-1}\mathbb{E}(xy) = \beta.$$

by SLLN and CMT.

Asymptotic Normality

Assume $\mathrm{Var}(xu) = \mathbb{E}(u^2 xx')$ exists. Then:

$$\sqrt{n}(\hat{\beta}_n - \beta) \xrightarrow{d} N(0, \Sigma),$$

where $\Sigma = \mathbb{E}(xx')^{-1}\mathbb{E}(u^2 xx')\mathbb{E}(xx')^{-1}$.
Proof: $\sqrt{n}(\hat{\beta}_n - \beta) = \left(\frac{1}{n}\sum x_i x_i'\right)^{-1}\frac{1}{\sqrt{n}}\sum x_i u_i$. CLT gives $\frac{1}{\sqrt{n}}\sum x_i u_i \xrightarrow{d} N(0, \mathbb{E}(u^2 xx'))$, then apply Slutsky.

Variance Estimation: Homoskedastic Case

Under $\mathbb{E}(u|x) = 0$, $\mathrm{Var}(u|x) = \sigma^2$: $\Sigma = \sigma^2\mathbb{E}(xx')^{-1}$. Estimate:

$$\hat{\Sigma} = \hat{\sigma}^2\left(\frac{1}{n}\sum x_i x_i'\right)^{-1}, \quad \hat{\sigma}^2 = \frac{1}{n}\sum \hat{u}_i^2.$$

Variance Estimation: Heteroskedastic Case

Without homoskedasticity, use the **Eicker-Huber-White** (robust) estimator:

$$\hat{\Sigma} = \left(\frac{1}{n}\sum x_i x_i'\right)^{-1}\left(\frac{1}{n}\sum \hat{u}_i^2 x_i x_i'\right)\left(\frac{1}{n}\sum x_i x_i'\right)^{-1}.$$

Consistency of $\hat{\Sigma}$ (Key Proof Step)

Need $\frac{1}{n}\sum \hat{u}_i^2 x_i x_i' \xrightarrow{p} \mathbb{E}(u^2 xx')$. Decompose:

$$\frac{1}{n}\sum \hat{u}_i^2 x_i x_i' = \frac{1}{n}\sum u_i^2 x_i x_i' + \frac{1}{n}\sum (\hat{u}_i^2 - u_i^2) x_i x_i'.$$

First term $\xrightarrow{a.s.} \mathbb{E}(u^2 xx')$ by SLLN. Second term $= o_p(1)$ because:

$$|\hat{u}_i^2 - u_i^2| = |x_i'(\beta - \hat{\beta}_n)| \cdot |\hat{u}_i + u_i|$$
$$\max_{i\le n} |\hat{u}_i^2 - u_i^2| \le \|\beta - \hat{\beta}_n\|^2 \max \|x_i\|^2 + 2\|\beta - \hat{\beta}_n\| \max \|x_i u_i\|.$$

Use the lemma: if $\mathbb{E}(\|Z_i\|^r) < \infty$ and Z_i identically distributed, then $\frac{\max_{i\le n}\|Z_i\|}{n^{1/r}} \xrightarrow{p} 0$.

Applied: $\frac{\max\|x_i\|^2}{n} = o_p(1)$, $\frac{\max\|x_i u_i\|}{\sqrt{n}} = o_p(1)$, and $\sqrt{n}(\hat{\beta} - \beta) = O_p(1)$.

Hypothesis Testing

Definitions

Null: $H_0 : \theta_0 \in \Theta_0$. Simple if Θ_0 singleton; composite otherwise.
Test: $\phi_n(X_1, \dots, X_n) \to \{0, 1\}$. Reject H_0 iff $\phi_n = 1$.
Type I Error: Reject when H_0 true. **Type II:** Fail to reject when false.
Power function: $\beta_n(\theta) = P_\theta(\phi_n = 1)$.
Size: $\alpha := \sup_{\theta\in\Theta_0} \beta_n(\theta)$.
Asymptotic size α**:** $\lim_{n\to\infty} \sup_{\theta\in\Theta_0} \beta_n(\theta) \le \alpha$.

Confidence Sets

C_n is a $1-\alpha$ confidence set if $P_\theta(\theta \in C_n) \ge 1 - \alpha$ for all θ.
Pivot: A function of data and unknown parameters whose distribution doesn't depend on unknown parameters (e.g. $\frac{\sqrt{n}(\bar{X}_n - \mu)}{\sigma} \sim N(0, 1)$).
Exact (known σ^2): $C_n = \left[\bar{X}_n \pm \frac{\sigma}{\sqrt{n}} z_{1-\alpha/2}\right]$.
Exact (unknown σ^2, normal): $\frac{\sqrt{n}(\bar{X}_n - \mu)}{\tilde{S}_n} \sim t_{n-1}$, yielding $C_n = \left[\bar{X}_n \pm \frac{\tilde{S}_n}{\sqrt{n}} t_{n-1,1-\alpha/2}\right]$, where $\tilde{S}_n^2 = \frac{1}{n-1}\sum (X_i - \bar{X}_n)^2$.

Finite Sample Inference (Normal Regression)

Under $Y|X \sim N(X\beta, \sigma^2 I_n)$: $\hat{\beta}|X \sim N(\beta, \sigma^2 (X'X)^{-1})$, $\frac{(n-k-1)\hat{\sigma}^2}{\sigma^2} \sim \chi^2_{n-k-1}$, and $\hat{\beta} \perp \hat{\sigma}^2 | X$.
t-statistic: $T_n = \frac{\hat{\beta}_j - \beta_{j,0}}{se(\hat{\beta}_j)} \sim t_{n-k-1}$, where $se(\hat{\beta}_j) = \hat{\sigma}\sqrt{e_j'(X'X)^{-1}e_j}$.
Reject $H_0 : \beta_j = \beta_{j,0}$ if $|T_n| > t_{n-k-1,1-\alpha/2}$.
p-value: $\hat{p} = 2F(-|T_n|)$ where F is the t_{n-k-1} CDF.

Testing Single Linear Restriction (Asymptotic)

$H_0 : r'\beta = c$. Under $\sqrt{n}(\hat{\beta}_n - \beta) \xrightarrow{d} N(0, V)$ with $\hat{V}_n \xrightarrow{p} V$:

$$T_n = \frac{\sqrt{n}(r'\hat{\beta}_n - c)}{\sqrt{r'\hat{V}_n r}} \xrightarrow{d} N(0, 1) \text{ under } H_0.$$

Reject if $|T_n| > z_{1-\alpha/2}$. CI: $C_n = r'\hat{\beta}_n \pm z_{1-\alpha/2}\sqrt{r'\hat{V}_n r/n}$.

Testing Multiple Linear Restrictions

$H_0 : R\beta = c$, R is $p \times (k + 1)$ full row rank.

$$T_n = n \cdot (R\hat{\beta}_n - c)'(R\hat{V}_n R')^{-1}(R\hat{\beta}_n - c) \xrightarrow{d} \chi^2_p.$$

Reject if $T_n > \chi^2_{p,1-\alpha}$. Confidence set: ellipsoid $\{c : T_n(c) \le \chi^2_{p,1-\alpha}\}$.
RVR' positive definite because: if $a \ne 0$, $R'a \ne 0$ (full rank), so $(R'a)'V(R'a) > 0$.

Testing Non-Linear Restrictions

$H_0 : f(\beta) = 0$, $f : \mathbb{R}^{k+1} \to \mathbb{R}^p$ continuously differentiable. Delta method:

$$\sqrt{n}(f(\hat{\beta}_n) - f(\beta)) \xrightarrow{d} N(0, D_\beta f(\beta)\, V\, D_\beta f(\beta)').$$

Construct χ^2_p statistic as before. Note $f(\beta) = R\beta$ yields linear case since $D_\beta f = R$.

Potential Outcomes & Causality

Setup

Individual i has potential outcomes $y_i(1)$ (treated) and $y_i(0)$ (untreated). Treatment $D_i \in \{0, 1\}$. Observed outcome:

$$Y_i = y_i(1)D_i + y_i(0)(1 - D_i).$$

Fundamental problem: never observe both $y_i(1)$ and $y_i(0)$.

Treatment Effects

ATE: $\mathbb{E}(y(1) - y(0))$.
ATT: $\mathbb{E}(y(1) - y(0)|D = 1)$.
ATU: $\mathbb{E}(y(1) - y(0)|D = 0)$.
Decomposition:

$$\text{ATE} = \text{ATT} \cdot P(D = 1) + \text{ATU} \cdot P(D = 0).$$

Naive Comparison and Selection Bias

$$\begin{aligned}&\mathbb{E}(Y|D = 1) - \mathbb{E}(Y|D = 0)\\ &= \underbrace{\mathbb{E}(y(1) - y(0)|D = 1)}_{\text{ATT}} + \underbrace{\mathbb{E}(y(0)|D = 1) - \mathbb{E}(y(0)|D = 0)}_{\text{Selection Bias}}.\end{aligned}$$

Naive comparison = ATT only if selection bias = 0.

Random Assignment

$D \perp (y(0), y(1))$ implies:

$$\mathbb{E}(y(d)|D) = \mathbb{E}(y(d)) \quad \text{for } d \in \{0, 1\}.$$

Selection bias vanishes, and $\beta_1 = \mathbb{E}(Y|D = 1) - \mathbb{E}(Y|D = 0) = \text{ATE}$. OLS of Y on D gives unbiased estimate of ATE.

Conditional Independence (Unconfoundedness)

$y(0), y(1) \perp D|w$ (selection on observables). Then:

$$\text{ATE} = \mathbb{E}[\mathbb{E}(Y|D = 1, w) - \mathbb{E}(Y|D = 0, w)].$$

Requires **overlap**: $0 < P(D = 1|w = w') < 1$ for all w'.

Homogeneous vs. Heterogeneous Effects

Homogeneous: $y_i(1) - y_i(0) = \beta_1$ for all i. Then $y_i = \beta_0 + \beta_1 D_i + u_i$ has a causal interpretation: β_1 is the treatment effect.
Heterogeneous: Effects vary across i. Regression coefficient is an average effect, not the individual effect.

Heterogeneous Effects with Interactions

If $x \in \{0, 1\}$ and effects vary, the correct specification is:

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 D_i + \beta_3 D_i x_i + v_i, \;\; \mathbb{E}(v|D, x) = 0.$$

Here $\beta_2 = \mathbb{E}(y(1)-y(0)|x=0)$ and $\beta_3 = \text{ATE}(x=1) - \text{ATE}(x=0)$.
Misspecification trap: If you omit $D_i x_i$ and run $y = b_0 + b_1 x + b_2 D + e$, then b_2 converges to a *variance-weighted* average of conditional ATEs:

$$b_2 \xrightarrow{p} \sum_x \frac{\text{Var}(D|x)\, P(x)}{\mathbb{E}(\text{Var}(D|x))} \cdot \text{ATE}(x),$$

which generally $\ne$ ATE unless effects are homogeneous or $P(D=1|x)$ is constant.

Inverse Probability Weighting (IPW)

Under unconfoundedness and overlap, with $p(x) := P(D=1|X=x)$:

$$\text{ATE} = \mathbb{E}\left[\frac{Y(D - p(X))}{p(X)(1 - p(X))}\right],$$
$$\text{ATT} = \mathbb{E}\left[\frac{YD}{P(D=1)}\right] - \mathbb{E}\left[\frac{Y(1-D)p(X)}{P(D=1)(1-p(X))}\right].$$

Useful when the propensity score $p(x)$ is easier to model than $\mathbb{E}(Y|D, X)$.

Multiple Treatments

k treatments, $k + 1$ potential outcomes. With random assignment:

$$\mathbb{E}(Y|x) = \beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k,$$

where $\beta_0 = \mathbb{E}(y(0))$ and $\beta_j = \mathbb{E}(y(j) - y(0))$.

Instrumental Variables

The Endogeneity Problem

If $\mathbb{E}(xu) \ne 0$ (endogeneity), OLS is inconsistent:

$$\hat{\beta}_n^{\text{OLS}} \xrightarrow{p} \beta + \mathbb{E}(xx')^{-1}\mathbb{E}(xu) \ne \beta.$$

IV Conditions

Use instrument $z \in \mathbb{R}^{l+1}$ ($z_0 = 1$) satisfying:
Validity: $\mathbb{E}(zu) = 0$ (exogeneity + exclusion).
Relevance (Rank Condition): $\mathbb{E}(zx')$ has rank $k + 1$.
Order Condition (necessary): $l \ge k$ (at least as many instruments as regressors).
$l = k$: exactly identified. $l > k$: overidentified.

Identification

From $y = x'\beta + u$ and $\mathbb{E}(zu) = 0$: $\mathbb{E}(zy) = \mathbb{E}(zx')\beta$. If $l = k$:

$$\beta = \mathbb{E}(zx')^{-1}\mathbb{E}(zy).$$

IV Estimator (Exact Identification)

$$\hat{\beta}_{\text{IV}} = \left(\frac{1}{n}\sum z_i x_i'\right)^{-1} \frac{1}{n}\sum z_i y_i = (Z'X)^{-1}Z'Y.$$

Consistent by LLN + CMT.

Potential Outcomes Framework for IV

Exclusion: $y(d,z') = y(d,z'')\ \forall d, z', z''$. Write $y(d) \equiv y(d,z')$.
Exogeneity: $y(d,z') \perp z|w\ \forall d, z'$.
Under constant linear treatment effects: $y(d) = x(w,d)'\beta + u$, $\mathbb{E}(u|w) = 0$.
Exclusion $\implies \mathbb{E}(u|z,w) = 0$, so $\mathbb{E}(zu) = 0$.

Asymptotic Distribution

$$\sqrt{n}(\hat{\beta}_{\text{IV}} - \beta) \xrightarrow{d} N(0, V_{\text{IV}}),$$

where in the scalar case with $\mathbb{E}(u^2|z) = \sigma^2$: $V_{\text{IV}} = \frac{\sigma^2}{\text{Corr}(x,z)^2\text{Var}(x)}$.

IV vs. OLS efficiency: If $\mathbb{E}(u|x) = 0$, OLS more efficient: $V_{\text{OLS}} = \frac{\sigma^2}{\text{Var}(x)}$.

GMM / 2SLS (Overidentification)

When $l > k$: $\mathbb{E}(zx')$ is $l+1 \times k+1$, not square. Use GMM:

$$\hat{\beta}_{\text{2SLS}} = (X'P_Z X)^{-1}X'P_Z Y,$$

where $P_Z = Z(Z'Z)^{-1}Z'$. Equivalently: regress X on Z (first stage), then Y on $\hat{X}$ (second stage).

Overidentification Tests (Sargan/Hansen)

Test $H_0 : \mathbb{E}(zu) = 0$. Under homoskedasticity, the F-statistic from regressing $\hat{u}$ on z is asymptotically $\chi^2_{l_2-k_1}$. Rejection $\implies$ invalid instruments or misspecification, but cannot tell which instrument is bad.

Must Include Exogenous Regressors in First Stage

With $x = (x_{-k}, x_k)$, endogenous x_k, instruments $z = (x_{-k}, z_k)$: the first stage *must* regress x_k on *all* of z, not just z_k.
Why: If you omit x_{-k}, the first-stage estimate $\tilde{\pi}$ suffers OVB: $\tilde{\pi} \xrightarrow{p} \pi + \frac{\text{Cov}(z_k, x_{-k})}{\text{Var}(z_k)}\delta \neq \pi$. The 2SLS estimate becomes $\hat{\beta}_k = \hat{\gamma}_k/\tilde{\pi} \xrightarrow{p} \beta_k\pi/(\pi+\nu) \neq \beta_k$, where $\hat{\gamma}_k$ is the reduced-form coefficient.
Intuitively, $\hat{x}_k$ without controls captures variation from x_{-k}, violating ceteris paribus.

IV Finite Sample Bias

IV is generally biased in finite samples. With one excluded instrument and one endogenous variable, $\hat{\beta}_{\text{IV}}$ has *no finite moments.* For proof: $\mathbb{E}[\hat{\beta}_{\text{IV}}|Z] = \beta + \mathbb{E}\left[\frac{\sum z_i u_i}{\sum z_i x_i}\middle| Z\right]$, and the ratio need not have finite expectation.

Hausman Test (Exogeneity)

Test $H_0 : \mathbb{E}(xu) = 0$ using both OLS and IV. Under H_0, both are consistent; under H_1, only IV is. Joint distribution:

$$\sqrt{n}\begin{pmatrix}\hat{\beta}_{\text{OLS}} - \beta \\ \hat{\beta}_{\text{IV}} - \beta\end{pmatrix} \xrightarrow{d} N(0, V_{\text{joint}}).$$

Test statistic: $T_n = n(\hat{\beta}_{\text{IV}} - \hat{\beta}_{\text{OLS}})'\hat{W}^{-1}(\hat{\beta}_{\text{IV}} - \hat{\beta}_{\text{OLS}}) \xrightarrow{d} \chi^2_k$ under H_0, where $\hat{W}$ consistently estimates the variance of the difference.

Weak Instruments

Setup

$y = \beta x + u$, $x = \pi z + v$, $\mathbb{E}(zu) = \mathbb{E}(zv) = 0$. Identification requires $\pi \neq 0$.

The Problem with $\pi \approx 0$

$$\hat{\beta}_{\text{IV}} = \beta + \frac{\frac{1}{\sqrt{n}}\sum z_i u_i}{\frac{1}{n}\sum z_i x_i}.$$

If $\pi = 0$: denominator $\xrightarrow{p} 0$ but numerator $\xrightarrow{d}$ normal. Can't apply Slutsky. The ratio converges to a **ratio of correlated normals**, not $N(0,V)$.

Finite Sample Bias

With joint normal errors and fixed instruments, $\hat{\beta}_{\text{IV}}$ has number of moments equal to (# excluded instruments − # endogenous variables). With $K = 1$ excluded instrument and 1 endogenous variable: $\hat{\beta}_{\text{IV}}$ has **no finite moments**.
$\hat{\beta}_{\text{OLS}} \xrightarrow{p} \frac{\sigma_{uv}}{\sigma_v^2}$ (biased toward OLS probability limit).

Rule of Thumb

First-stage F-statistic ≥ 10 for relative bias $\leq 10\%$ (Stock & Yogo, 2005). Critical values range 9–12 for 3–30 instruments. But this is only an approximation.

Anderson-Rubin Test

Robust to weak instruments. Suppose $y = x'\beta + u$, $\mathbb{E}(zu) = 0$.
Under $H_0 : \beta = \beta_0$: $\mathbb{E}(z(y - x'\beta_0)) = 0$, so regress $u(\beta_0) = y - x'\beta_0$ on z:

$$u_i(\beta_0) = z_i'\gamma + \epsilon_i, \quad \mathbb{E}(z\epsilon) = 0.$$

Under H_0: $\gamma = 0$. Test statistic:

$$T_n = n\hat{\gamma}'\hat{V}^{-1}\hat{\gamma} \xrightarrow{d} \chi^2_l.$$

Reject if $T_n > \chi^2_{l,1-\alpha}$.

Why robust: Under H_0, $\sqrt{n}\hat{\gamma} = \left(\frac{1}{n}\sum z_i z_i'\right)^{-1}\frac{1}{\sqrt{n}}\sum z_i u_i$. This uses CLT on $z_i u_i$ directly—no division by a possibly-near-zero first stage.
Power: Under $H_1 : \beta \neq \beta_0$, we have $\gamma = \mathbb{E}(zz')^{-1}\mathbb{E}(zx')(\beta - \beta_0) \neq 0$, so the test has power against any alternative.

With Included Instruments

Separate $z = (z_1, z_2)$, $x = (x_1, z_1)$. Under $H_0 : \beta_1 = \beta_{1,0}$, regress $y - x_1'\beta_{1,0}$ on z_1 and z_2. Test whether coefficients on z_2 are zero:

$$T_n = n\hat{\gamma}_{z_2}'\hat{V}^{-1}\hat{\gamma}_{z_2} \xrightarrow{d} \chi^2_{l_2}.$$

Difference-in-Differences

Setup

Two periods: $T \in \{0,1\}$. Two groups: $G \in \{0,1\}$ (treated/control). Observed outcome:

$$Y = \begin{cases} y(1) & \text{if } G = 1, T = 1 \\ y(0) & \text{otherwise.}\end{cases}$$

Target: $\text{ATT} = \mathbb{E}(y(1) - y(0)|G = 1, T = 1)$.

Naive Comparisons Fail

Across time (treated group):

$$\mathbb{E}(Y|G{=}1, T{=}1) - \mathbb{E}(Y|G{=}1, T{=}0) = \text{ATT} + \text{Temporal trend.}$$

Across groups (post-period):

$$\mathbb{E}(Y|G{=}1, T{=}1) - \mathbb{E}(Y|G{=}0, T{=}1) = \text{ATT} + \text{Selection bias in } y(0).$$

Common Trends Assumption

In the absence of treatment, the change in outcomes would be the same for treated and control groups:

$$\begin{aligned}&\mathbb{E}(y(0)|G = 1, T = 1) - \mathbb{E}(y(0)|G = 1, T = 0)\\ &= \mathbb{E}(y(0)|G = 0, T = 1) - \mathbb{E}(y(0)|G = 0, T = 0).\end{aligned}$$

Note: the LHS involves an *unobservable* counterfactual.

DiD Estimand

Under common trends:

$$\begin{aligned}\text{ATT} = &[\mathbb{E}(Y|G{=}1, T{=}1) - \mathbb{E}(Y|G{=}1, T{=}0)]\\ &- [\mathbb{E}(Y|G{=}0, T{=}1) - \mathbb{E}(Y|G{=}0, T{=}0)].\end{aligned}$$

Regression Implementation

$$Y_{it} = \beta_0 + \beta_1 G_i + \beta_2 T_t + \beta_3(G_i \cdot T_t) + u_{it}.$$

β_3 is the DiD estimator = ATT under common trends.
β_1: group difference at baseline. β_2: common time effect.

Data Requirements

Repeated cross section: Random sample in each period (different units).
Panel data: Same units observed in both periods (stronger).
Every panel is a repeated cross section, but not vice versa.

Regression Discontinuity Design

Setup

Running variable X, cutoff c, treatment D. Potential outcomes $y(0), y(1)$.

Sharp RDD

$D = \mathbf{1}(X \geq c)$: treatment is deterministic in X.
Unconfoundedness holds: $y(0), y(1) \perp D|X$ (since D is a function of X).
Overlap fails: $P(D=1|X=x) = \mathbf{1}(x \geq c) \in \{0,1\}$.

Identification

Under continuity: $\mathbb{E}(y(0)|X=x)$ continuous at c:

$$\mathbb{E}(y(1)-y(0)|X=c) = \mathbb{E}(Y|X=c) - \lim_{x\uparrow c}\mathbb{E}(Y|X=x).$$

Identifies treatment effect **at the cutoff only**.

Estimation: Local Linear Regression

Choose bandwidth h and solve:

$$\min_{\alpha_0,\beta_0,\gamma,\delta}\sum_{i=1}^{n}\mathbf{1}(|X_i-c|\leq h)(y_i-\alpha_0-\beta_0 x_i-\gamma d_i-\delta d_i x_i)^2.$$

Equivalent to two separate regressions on $\{i : X_i \in [c-h, c)\}$ and $\{i : X_i \in [c, c+h]\}$.
Recentering: $Y_i = \alpha_0 + \beta_0(X_i - c) + \gamma D_i + \delta D_i(X_i - c) + \epsilon_i$ makes γ the discontinuity.

Bandwidth Choice

Optimal: $h = C \cdot n^{-1/5}$ (bias-variance tradeoff). Larger $h \Rightarrow$ lower variance, higher bias.
IK (2012) and CCT (2014) propose data-driven bandwidth selectors. CCT accounts for asymptotic bias.

Threats to Validity

Manipulation: Individuals choosing X values near cutoff.
McCrary (2008): test for density discontinuity at c.
Multiple treatments: Cannot identify which treatment caused the jump.
Covariate balance: Pre-determined covariates should be continuous at c; discontinuities suggest violations.

Fuzzy RDD

$P(D=1|X=x)$ is discontinuous at c, but $D \neq \mathbf{1}(X \geq c)$. Then $Z = \mathbf{1}(X \geq c)$ is an instrument for D.
Under monotonicity ($P(D_1 \geq D_0) = 1$), the estimand is a LATE at the cutoff:

$$\begin{aligned}&\mathbb{E}(y(1)-y(0)|X=c,\text{complier})\\ &= \frac{\lim_{x\downarrow c}\mathbb{E}(Y|X=x)-\lim_{x\uparrow c}\mathbb{E}(Y|X=x)}{\lim_{x\downarrow c}\mathbb{E}(D|X=x)-\lim_{x\uparrow c}\mathbb{E}(D|X=x)}.\end{aligned}$$

Fuzzy RDD Implementation

2SLS on subsample $\{i : |X_i - c| \leq h\}$:
First stage: $D = \pi_0 + \pi_1 Z + \pi_2(X-c) + \pi_3 Z(X-c) + v$.
Second stage: $Y = \beta_0 + \beta_1 D + \beta_2(X-c) + \beta_3 Z(X-c) + u$.

Panel Data

Setup

N individuals, T time periods. Linear model:

$$Y_{it} = X_{it}'\beta + \alpha_i + u_{it},$$

where α_i is an unobserved individual fixed effect.
Asymptotics: Large N, small T.

Problem with Pooled OLS

Pooled OLS treats $\alpha_i + u_{it}$ as the error. If $\mathbb{E}(X_{it}\alpha_i) \neq 0$, the composite error is correlated with regressors, and pooled OLS is inconsistent.

First Differencing (FD)

Difference across time to eliminate α_i:

$$\Delta Y_{it} = \Delta X_{it}'\beta + \Delta u_{it}.$$

FD estimator: $\hat\beta_{\text{FD}}$ is OLS applied to differenced data.
Consistency requires: $\mathbb{E}(\Delta X_{it}'\Delta u_{it}) = 0$, i.e. **contemporaneous exogeneity**: $\mathbb{E}(u_{it}|X_{it}, X_{it-1}) = 0$.
Not sufficient if: unobservables in *other* time periods are correlated with today's regressors.

Strict Exogeneity

$\mathbb{E}(u_{it}|X_{i1},\ldots,X_{iT}) = 0$ for all t. Stronger than contemporaneous exogeneity. Required for FE consistency.
Violated if, e.g., past outcomes affect future regressors (feedback effects).

Fixed Effects (FE) / Within Estimator

For general $T \geq 2$, define within-transformed variables $\ddot{Y}_{it} = Y_{it} - \bar{Y}_i$:

$$\ddot{Y}_{it} = \ddot{X}_{it}'\beta + \ddot{u}_{it}.$$

FE estimator is OLS on demeaned data. Eliminates α_i without differencing.
Under strict exogeneity: $\hat\beta_{\text{FE}}$ is consistent as $N \to \infty$ (fixed T).
FE = LSDV equivalence: Regressing Y on X and N individual dummies (LSDV) yields the same $\hat\beta$ as the within estimator. **Proof:** By FWL, $\hat\beta_{\text{LSDV}} = (X'M_D X)^{-1}X'M_D Y$ where D is the matrix of individual dummies. M_D demeans within each individual: $(M_D Y)_{it} = Y_{it} - \bar{Y}_i = \ddot{Y}_{it}$. So $\hat\beta_{\text{LSDV}} = (\ddot{X}'\ddot{X})^{-1}\ddot{X}'\ddot{Y} = \hat\beta_{\text{FE}}$.

FD vs. FE

With $T = 2$: FD = FE.
With $T > 2$: differ in general. FE more efficient under homoskedasticity of u_{it}; FD more robust to serial correlation patterns.

Serial Correlation

Panel structure allows arbitrary correlation of u_{it} with α_i (eliminated by FD/FE), but standard errors must account for within-individual serial correlation in u_{it}. Cluster standard errors at the individual level.

Tips and Tricks

Proof Strategies for Consistency

1. Write estimator as function of sample averages.
2. Apply SLLN to each sample average.
3. Apply CMT to the composed function.

For extremum estimators (MLE, GMM): show uniform convergence of objective function + identification at θ_0.

Trace Trick for Quadratic Forms

For scalar $U'AU$: $U'AU = \text{tr}(U'AU) = \text{tr}(AUU')$, so $\mathbb{E}[U'AU|X] = \text{tr}(A\,\mathbb{E}[UU'|X])$. Key identity: $\text{tr}(AB) = \text{tr}(BA)$.
Used to prove $\mathbb{E}[\text{SSR}|X] = \sigma^2(n-k-1)$.

Proof Strategies for Asymptotic Normality

1. Decompose $\sqrt{n}(\hat\theta - \theta)$ into a CLT term and remainder.
2. Apply CLT to iid mean-zero term.
3. Show remainder is $o_p(1)$ using Slutsky.

Template: $\sqrt{n}(\hat\beta - \beta) = \left(\frac{1}{n}\sum x_i x_i'\right)^{-1}\frac{1}{\sqrt{n}}\sum x_i u_i$. CLT on second factor, Slutsky with first.

Useful Inequalities

Markov: $P(|X| \geq t) \leq \frac{\mathbb{E}|X|}{t}$.
Chebyshev: $P(|X-\mu| \geq t) \leq \frac{\text{Var}(X)}{t^2}$.
Jensen: If g convex, $g(\mathbb{E}(X)) \leq \mathbb{E}(g(X))$. Strict if g strictly convex and X non-degenerate.
Cauchy-Schwarz: $|\mathbb{E}(XY)|^2 \leq \mathbb{E}(X^2)\mathbb{E}(Y^2)$.

Key o_p/O_p Arguments

To show $\frac{1}{n}\sum \hat{u}_i^2 x_i x_i' \xrightarrow{p} \mathbb{E}(u^2xx')$:
Bound
$\max_{i\leq n}|\hat{u}_i^2 - u_i^2| \leq \|\hat\beta - \beta\|^2 \max\|x_i\|^2 + 2\|\hat\beta - \beta\|\max\|x_i u_i\|$.
Use: $\frac{\max\|Z_i\|}{n^{1/r}} = o_p(1)$ when $\mathbb{E}\|Z\|^r < \infty$.

Common Endogeneity Sources

- Omitted variables correlated with both x and y
- Simultaneity / reverse causality
- Measurement error in regressors
- Self-selection into treatment

IV Checklist

1. **Relevance:** First-stage F ≥ 10 (weak instrument check)
2. **Exclusion:** z affects y only through x (untestable)
3. **Exogeneity:** $\mathbb{E}(zu) = 0$ (partially testable via overid)
4. **Monotonicity:** For LATE interpretation with heterogeneous effects

Identification Strategy Summary

- **RCT:** Random assignment $\implies$ ATE from simple regression
- **Selection on observables:** Unconfoundedness + overlap $\implies$ ATE
- **IV:** Exogeneity + relevance $\implies$ causal effect (LATE if heterogeneous)
- **DiD:** Common trends $\implies$ ATT
- **RDD:** Continuity at cutoff $\implies$ treatment effect at cutoff
- **Panel FE/FD:** Eliminates time-invariant unobservables

Bias-Variance Tradeoff (Irrelevant Variables)

Including irrelevant variable ($\beta_2 = 0$): no bias, but *increases* variance. Omitting relevant variable ($\beta_2 \neq 0$): introduces OVB, but *decreases* variance. Via FWL: $\text{Var}(\hat\beta_1|X) = \sigma^2/\text{SSR}_1$ where SSR_1 is residual SS from regressing x_1 on other regressors. Adding correlated regressors lowers SSR_1, inflating variance.